Class activity – GAN

Name: Satyatma Chincholi

USN: 1RVU23CSE414

Date: 23/02/2026

Generative Adversarial Networks (GANs) have evolved rapidly from simple image generators to sophisticated tools capable of artistic style transfer and high-resolution synthesis.

1. DCGAN (Deep Convolutional GAN)

Description: This was the first architecture to successfully use **Convolutional Neural Networks (CNNs)** instead of fully connected layers. It introduced architectural constraints—like using strided convolutions instead of pooling and removing fully connected hidden layers—to make GAN training more stable.

- **Key Feature:** Uses "Fractionally-strided convolutions" (Upsampling) to generate images.
- **Task 1:** Take the MNIST dataset (handwritten digits) and build a DCGAN to generate new, synthetic digits.
- **Task 2:** Experiment with the **latent space**. Generate two different faces, then "average" their latent vectors to see if the resulting image looks like a blend of the two.

2. CGAN (Conditional GAN)

Description: Standard GANs are "unconditional," meaning you can't tell them *what* to generate; they just output a random sample from the training distribution. CGANs fix this by feeding a **label (y)** into both the Generator and Discriminator.

- **Key Feature:** Direct control over the class of the generated output (e.g., "generate a 7" or "generate a dog").
- **Task 1:** Modify your MNIST DCGAN to be "conditional." Use the digit labels (0-9) as input so you can request a specific number on demand.
- **Task 2:** Create a CGAN that generates clothing items based on a text tag like "Shirt" or "Shoe" using the Fashion-MNIST dataset.

3. Pix2Pix (Image-to-Image Translation)

Description: This is a type of Conditional GAN designed for **paired** image-to-image translation. It learns a mapping from an input image to an output image (e.g., a black-and-white sketch to a color photo).

- **Key Feature:** Uses a **U-Net** architecture for the Generator and a "PatchGAN" Discriminator that looks at local image patches rather than the whole image.
- **Task 1:** Use the "Edges-to-Shoes" dataset to train a model that turns a simple line drawing into a realistic-looking photo of a sneaker.
- **Task 2:** Train a model to convert daytime street photos into nighttime versions using a paired dataset of cityscapes.

4. CycleGAN (Unpaired Image Translation)

Description: Unlike Pix2Pix, CycleGAN does not require paired examples. It can learn to turn photos of horses into zebras even if you don't have a photo of the *exact same* horse as a zebra. It uses **Cycle Consistency Loss** to ensure the translated image can be "translated back" to the original.

- **Key Feature:** Works on two collections of images (Set A and Set B) without one-to-one mapping.
- **Task 1:** Perform a "Style Transfer" task, such as turning a standard landscape photo into the style of a Claude Monet painting.
- **Task 2:** Create a "Summer to Winter" filter that changes the foliage and lighting of a forest photo without changing the actual trees' positions.

5. StyleGAN

Description: Developed by NVIDIA, StyleGAN focuses on the **style** of the image at different scales (coarse features like pose, and fine features like hair color). It doesn't feed the latent code at the beginning; instead, it injects it at every layer of the generator.

- **Key Feature:** Progressive growing and "Adaptive Instance Normalization" (AdaIN) for incredible high-resolution detail.
- **Task 1:** Use a pre-trained StyleGAN (like StyleGAN2 or 3) to perform **"Face Swapping"** or attribute manipulation (e.g., adding glasses or changing the age of a generated face).
- **Task 2:** Explore the **Mapping Network**. Visualize how changing a single "style" parameter affects the output at different resolutions (e.g., changing skin texture vs. changing head shape).

Model	Input	Mapping Type	Best For
DCGAN	Random Noise	Unconditional	Basic image generation
CGAN	Noise + Label	Conditional	Specific category generation
Pix2Pix	Image A	Paired	Maps (satellite to street), Colorization
CycleGAN	Image A	Unpaired	Artistic style, Season changes
StyleGAN	Latent Code	Style-based	High-res photorealism (Faces)

Generating face using – “This person does not exist”